

# Testing in production

Losing fear, one test at a time

Jorge Marín

November 7th, 2019 – Bristech 

[@chipironcin](https://jorgemarin.me)

# Hidden slide

Welcome to Bristech and thanks for coming to my talk 🥳

I know we are in the second time, you are a bit tired already after attending to many interesting talks but I promise this talk won't be too heavy.

\*breathe\*


\*no, srsly, breathe\*

My name is Jorge, and I am an engineer 🖐️ 🇪🇸

# Testing in production

Losing fear, one test at a time

# Poll time - Raise your hands if you...


- Are having a good day
- Often forget to squash before merge
- Know what testing is
- Know about the testing pyramid or any other way of classifying different testing levels
- Have a production environment
- Are testing your services in production
- Prefer spaces to tabs 

# What is this talk about

- Mostly answers to Frequently Asked Questions
- Not going to be too technical
- Based on the experience gained working at Dyson

# Frequent Asked Questions

- Why do you need tests in production?
- What is the right testing level against production systems?
- How can we define an end to end test?
- How to avoid disrupting real users?
- How to keep your tests out of business reports?
- Cleanup
  - Should you clean after yourself? What if you don't?



Let's dive in!



Why do you need testing in production? #1




# What do you mean by “testing in production”?

lambda  
service

To guarantee that your **application** continues to work  
consistently after the release/deployment

# Why do you need testing in production? #1

- But I have the best real-time health dashboards in the world and will tell me when something is wrong



# Why do you need testing in production? #1

- Imagine a lambda function
  - One function that returns the current date
  - Unit tested
    - Upon call, returns a formatted string in the form “
  - Integration tested
    - API endpoint works as expected and other services are getting the string
  - Manually tested
  - No new versions, deploy, verify and forget
  - Dashboards are healthy
    - Service healthy
    - Metrics healthy
    - Lots of 200OK

# Why do you need testing in production? #1

- Imagine a lambda function
  - One function that returns the current date
  - Unit tested
    - Upon call, returns a formatted string in the form “
  - Integration tested
    - API endpoint works as expected and other services are getting the string
  - Manually tested
  - No new versions, deploy, verify and forget
  - Dashboards are healthy
    - Service healthy
    - Metrics healthy
    - Lots of 200OK

**UNTIL YEAR 2038**



# y2k38

- Storing dates (seconds since Jan 01 1970) as a signed 32-bit binary integer

32-bit binary representation	Seconds since Jan 01 1970	Date	
000000000000000000000000000000	0	Thursday, 01-Jan-1970 00:00:00 UTC	Beginning epoch
01011101110000111111011000000011	1573123587	Thursday, 07-Nov-2019 10:46:27 UTC	Today
011111111111111111111111111111	2147483647	Tuesday, 19-Jan-2038 03:14:07 UTC	y2k38
100000000000000000000000000000	-2147483648	Friday, 13-Dec-1901 20:45:52 UTC	One day after Guglielmo Marconi received the first trans-Atlantic radio signal, sent 1,700 miles from Poldhu in Cornwall, England to Signal Hill, St. John's in Newfoundland in Canada

y2k38



# Why do you need testing in production? #1

- Or just imagine one of your stateful services becomes inconsistent and starts failing for all customers
- At that moment, lots of customers will call you saying:
  - Cannot login
  - Cannot register
  - Cannot change profile picture
  - Cannot launch the app
  - Cannot order chocolate milkshakes takeaway


# Why do you need testing in production? #1

- What do you prefer:

Real angry **paying** customers calling you because what they paid for is not working




An automated process acting as the first real user noticing an issue and notifying the team right away







# What is the right testing level against production systems? #2




# What is the right testing level against production systems? #2




# What is the right testing level against production systems? #2

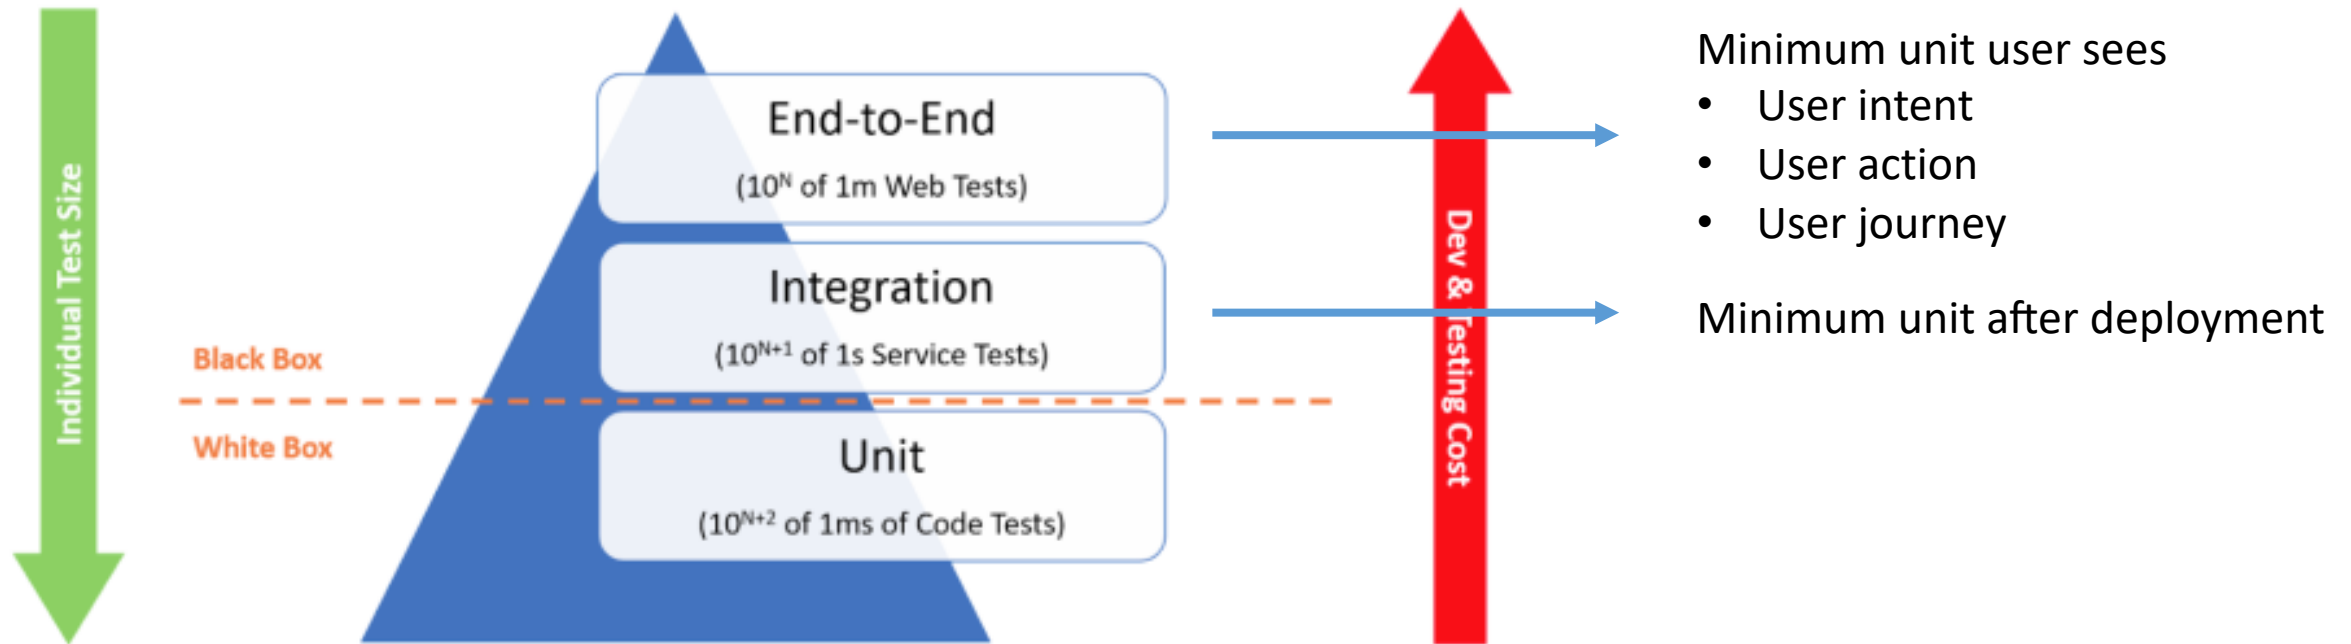


# What is the right testing level against production systems? #2






# What is the right testing level against production systems? #2



# How can we define an end to end test? #3

- More than 1 million connected machines performing automated actions like scheduled cleans
- Users all over the world
  - actions as simple as renaming your cleaning robot from 'Dyson 360 Eye' to 'RoboMop9000'
- Fast
- Easy to run
- Extendable
- Behavior focused




# How can we define an end to end test? #3



- Given an existing user with this configuration
- And the user logs in
- And the user changes robot name to 'RoboMop9000'
- Then the robot name has changed to 'RoboMop9000'



# How can we define an end to end test? #3




# How can we define an end to end test? #3

- Combination of:
  - CucumberJS
  - Request library
  - Assert library
  - Lots of single steps defining calls to APIs and data transformations
    - Mocking machine (virtual machine)
    - Mocking user (virtual user)
    - But going to the cloud and back
  - Grouped together under user journeys or scenarios



# How avoid disrupting real users? #4

- If your test is too close to real users and real actions...




# How avoid disrupting real users? #4

- Spacing these end to end tests
  - The more often you run them the faster you will notice an issue
  - Noise
  - Slowing down your production systems
  - Balance
- Virtual users registered in production like real users with their own configuration

# How to keep your tests out of statistics? #5

- But testing in production can also skew some of the metrics you record like:
  - Number of users interacting with your system
  - API usage
  - Error count
- Even mixed logs
  - How to discern what was part of a test or part of a real interaction?



**problem?**

# How to keep your tests out of statistics? #5

- Correlation id

ffffffff-0000-0000-0000-000000000000

- User agent


AppId/AppVersionId (Language=LanguageNameAndOptionallyVersion)

- Custom HTTP headers

TestToolVersion=1.1.2

# Cleanup? #6

- Running tests in production 👍



# Cleanup? #6

- But... what about the all the fake data they generated?





# Cleanup? #6

- Automatically clean as much testing items as possible after each test
- Objective: Leave the system as you found it, as if the tests never ran there



# Cleanup? #6

- Helps starting tests from a clean state
  - Avoiding false negatives
- Avoids cluttering your production databases
- Avoids exhausting other limited resources
- Avoids making real data more difficult to retrieve and search



# That's interesting, new questions!

- Can your E2E Tests pass and you still have problems?
- By their nature E2E Tests span multiple services and teams  
Who is responsible for them and how do teams collaborate to write them?
- How do E2E Tests compare to dark or canary Deployments?  
Do you need both?
- How do you avoid E2E Tests becoming fragile?



# Recap

- Testing in production is good and necessary 🎉
- Think like an user, act like an user 🧠
- CucumberJS is a good start
- Think about your system's capacity 📈
- Mark test intents to differentiate them from real user interactions
- Clear your test data and reset connections/status after each test

That's all folks



Before Q&A

**THE FINAL**

**BIT-BORING-TALK-BUT-WILL-SMILE-FOR-THE-GRAM**

**SELFIE WITH ATTENDEES!**

and you?  
are you brave enough to test in production?

Questions round